


For my performance analysis, I gathered data on the major page faults and wall-clock time returned by my proj6 executable when run on 3 files of varying sizes. The sizes of these files are 26 bytes, 6254 bytes, and 1253858 bytes. I ran the program with both standard I/O as well as memory-mapped I/O, and I also changed the chunk size for each of these runs. Here's a screenshot of how I've run my program on each of the files:

```
student@cs2324:~/CS502/proj6$ ./doit ./proj6 moby.txt 8192 mmap
File size: 1253858 bytes.
Strings of at least length 4: 24247
Maximum string length: 71 bytes


-Statistics-
*** System time = 0 milliseconds
*** User time = 3 milliseconds
*** Wall clock time = 3 milliseconds
*** Involuntary context switches = 0
*** Voluntary context switches = 2
*** Page faults requiring I/O = 0
*** Page faults serviced without I/O = 95
*** Maximum resident set size = 2272 kilobytes
```

The graphs of the standard I/O vs memory-mapped I/O results are shown below:

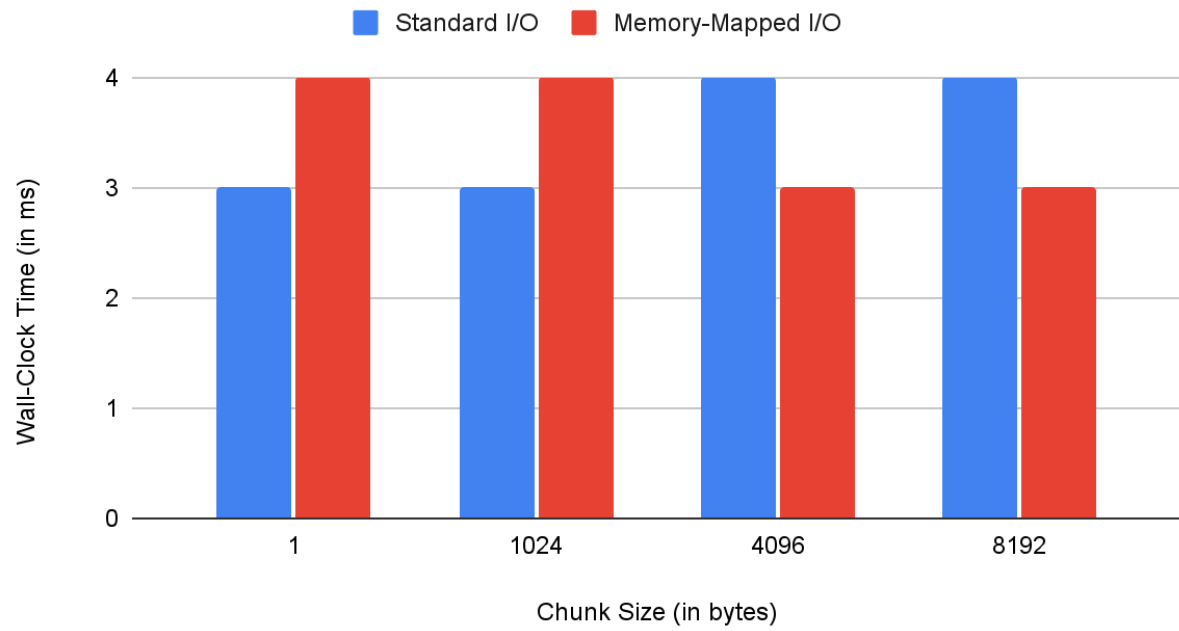
I/O for 26 Byte File



I/O for 6254 Byte File




I/O for 1 MB File




To understand the performance of the different thread options, I've only run the program on the 1 MB file since it provides the most meaningful results. I've pasted the graph of these results below:

Major Page Faults vs. Number of Threads



Wall-Clock Time (in ms) vs. Number of Threads



For all of my tests, my proj6 program never ran into any major page faults according to the statistics acquired from my doit program. The wall-clock time for both the 26 byte and 6254 byte files was always 0 or 1, which may be because of the severely small size of both files. As for the 1 MB file I tested with, the wall-clock time always hovered around 3 to 4 ms, and I think this may be because of the power of the system I am testing with. As for testing with the multi-threaded portion of my program, I wasn't able to gain meaningful results in this area either. Although I don't see much loss or improvement in performance for utilizing more or less threads in my tests, I believe that using higher thread counts for smaller files would probably increase the wall-clock time while higher thread counts for larger files would probably decrease the wall-clock time. My reasoning behind this is that the overhead required for more threads on a smaller file will probably decrease performance in that scenario, while increasing performance if used on a larger file.